

Introduction to the Course

INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON




Rob Reider

Adjunct Professor, NYU-Courant
Consultant, Quantopian

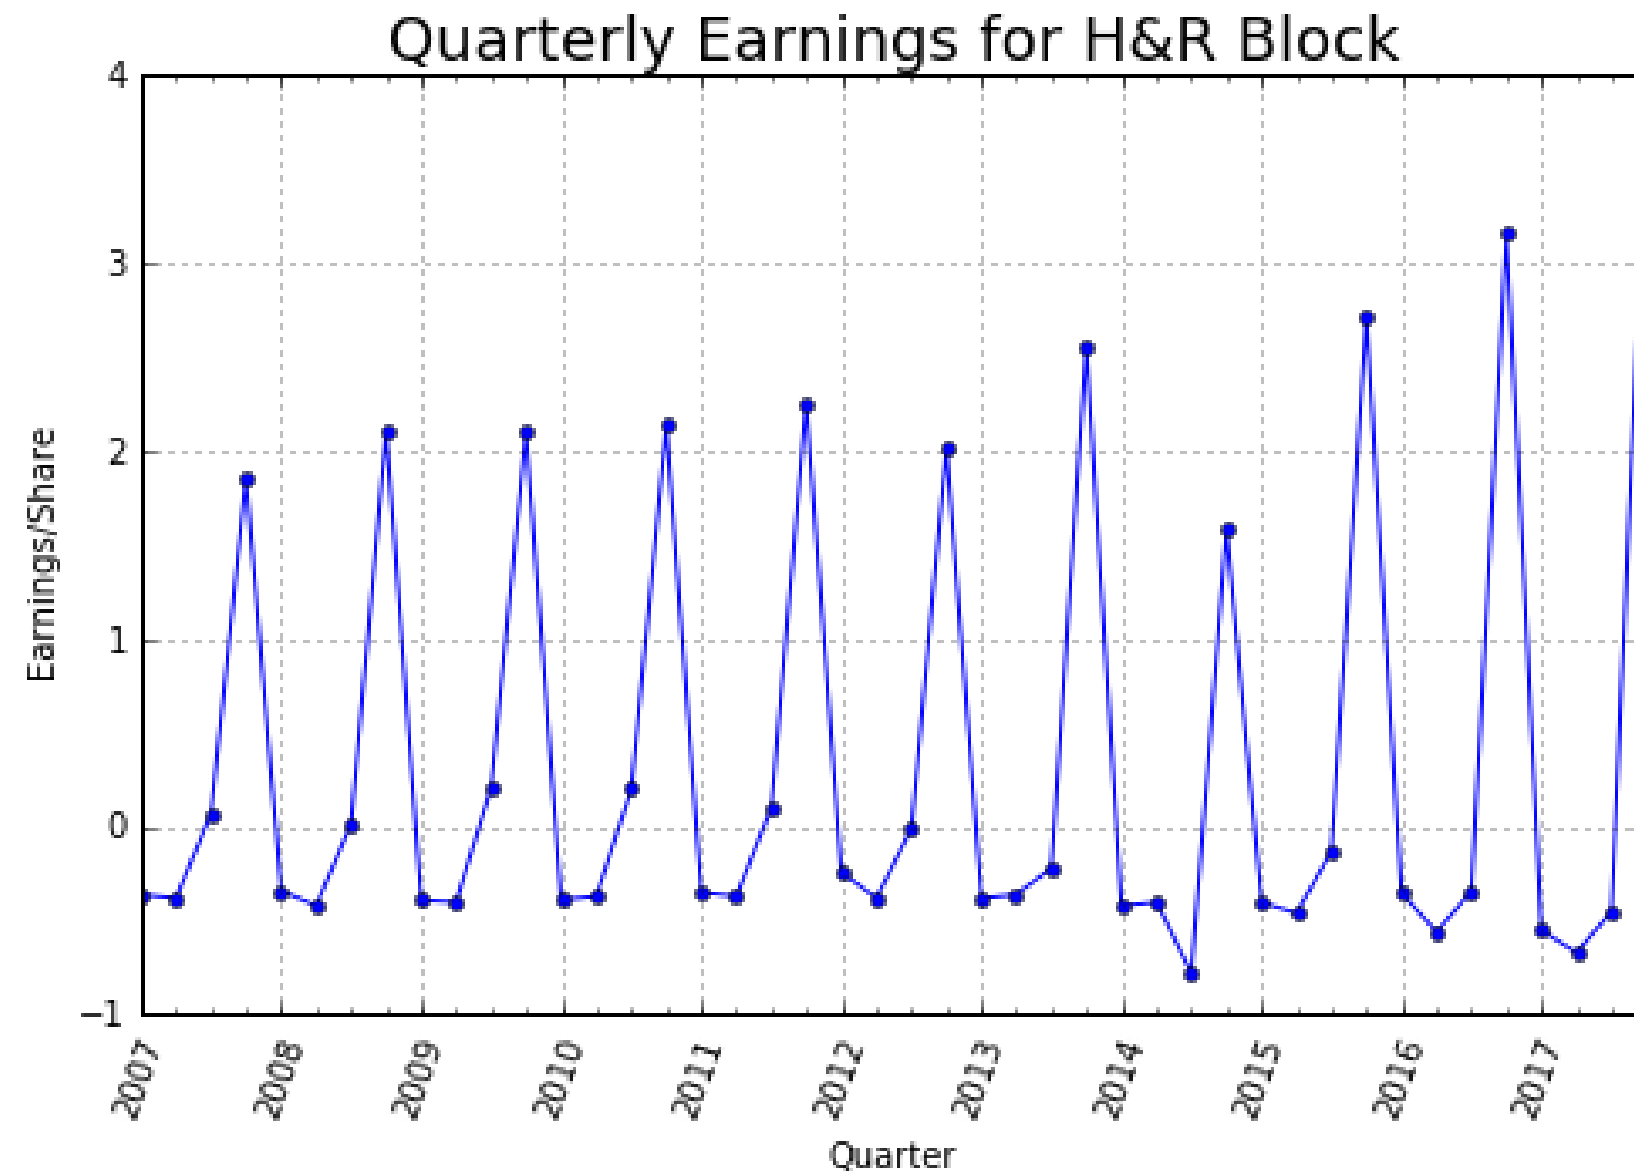
Example of Time Series: Google Trends




Example of Time Series: Climate Data



Example of Time Series: Quarterly Earnings Data



Example of Multiple Series: Natural Gas and Heating Oil



Goals of Course

- Learn about time series models
- Fit data to a times series model
- Use the models to make forecasts of the future
- Learn how to use the relevant statistical packages in Python
- Provide concrete examples of how these models are used

Some Useful Pandas Tools

- Changing an index to datetime

```
df.index = pd.to_datetime(df.index)
```

- Plotting data

```
df.plot()
```

- Slicing data

```
df[ '2012' ]
```

Some Useful Pandas Tools

- Join two DataFrames

```
df1.join(df2)
```

- Resample data (e.g. from daily to weekly)

```
df = df.resample(rule='W', how='last')
```


More pandas Functions

- Computing percent changes and differences of a time series

```
df['col'].pct_change()  
df['col'].diff()
```

- pandas correlation method of Series

```
df['ABC'].corr(df['XYZ'])
```

- pandas autocorrelation

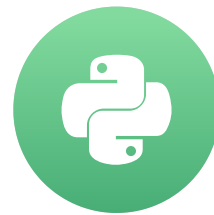
```
df['ABC'].autocorr()
```

Let's practice!

INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

Correlation of Two Time Series

INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

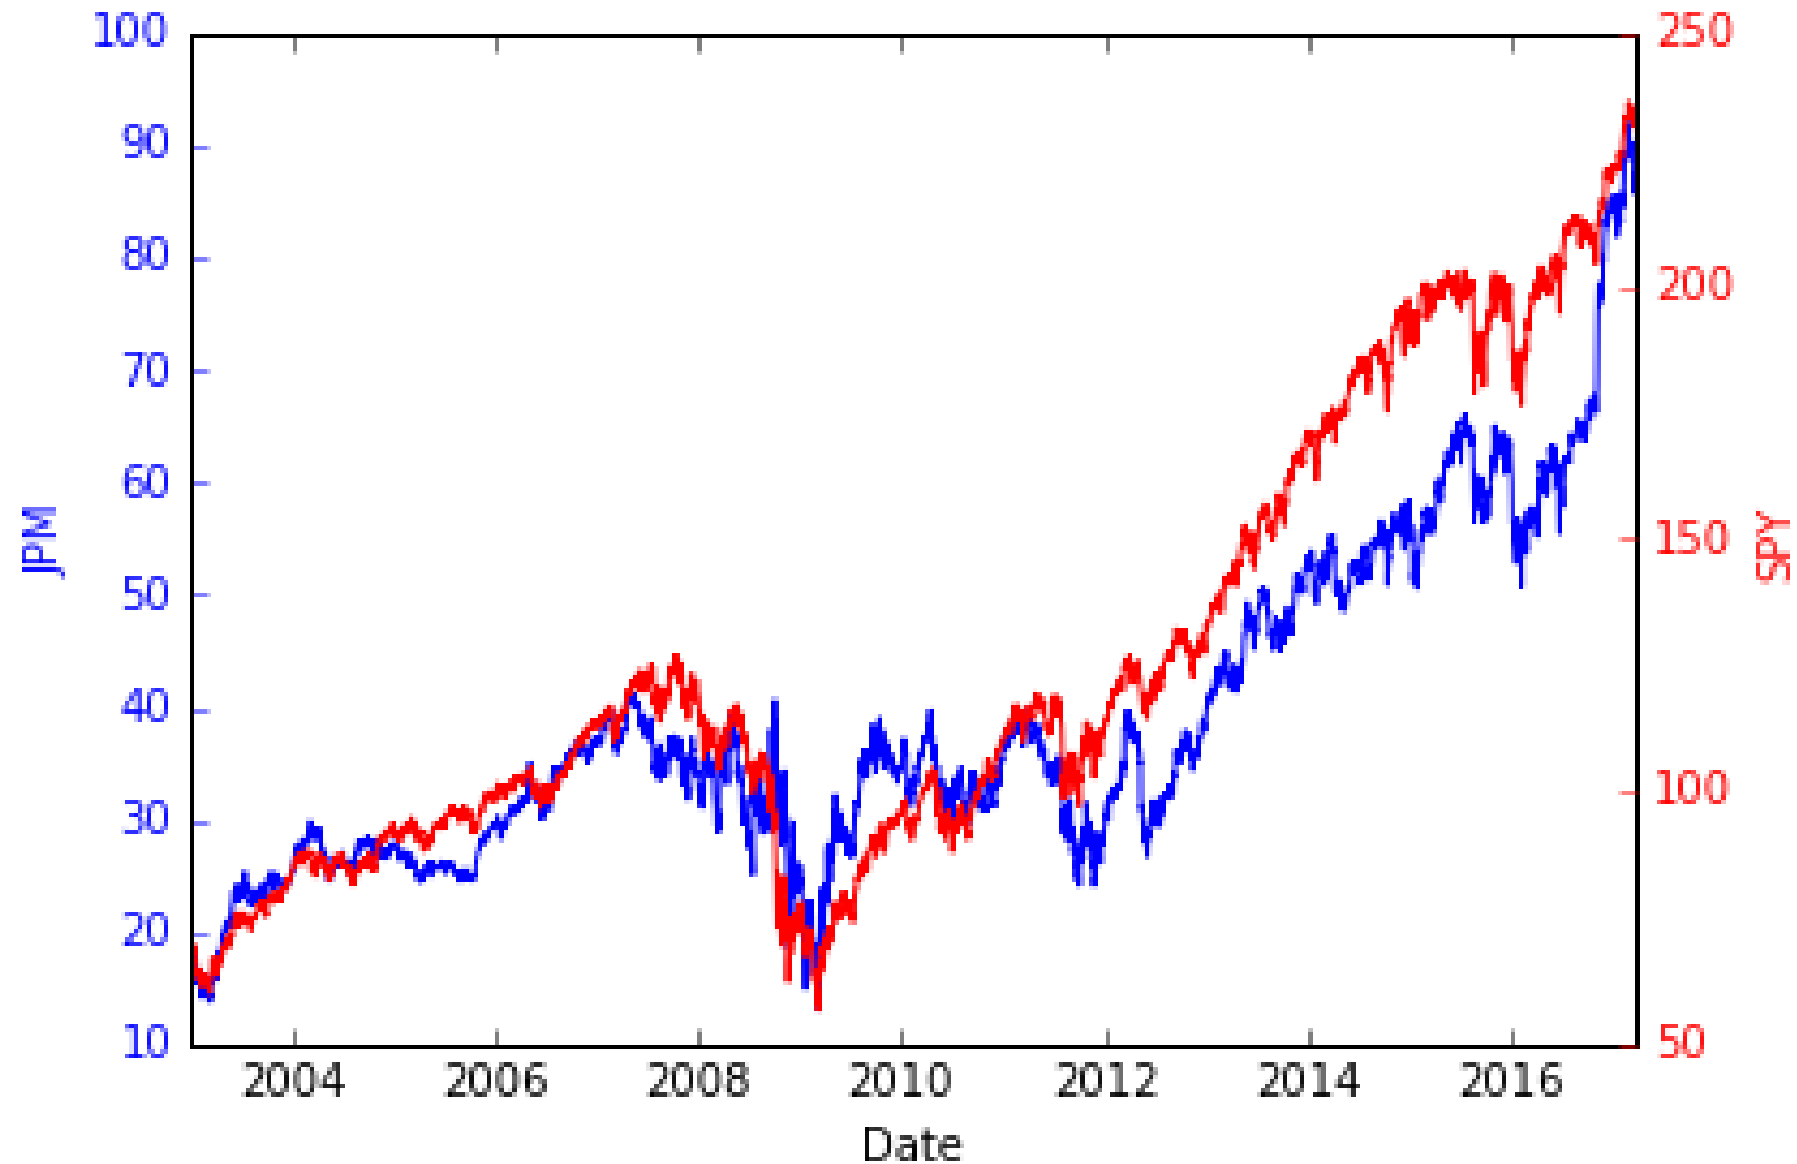


Rob Reider

Adjunct Professor, NYU-Courant
Consultant, Quantopian

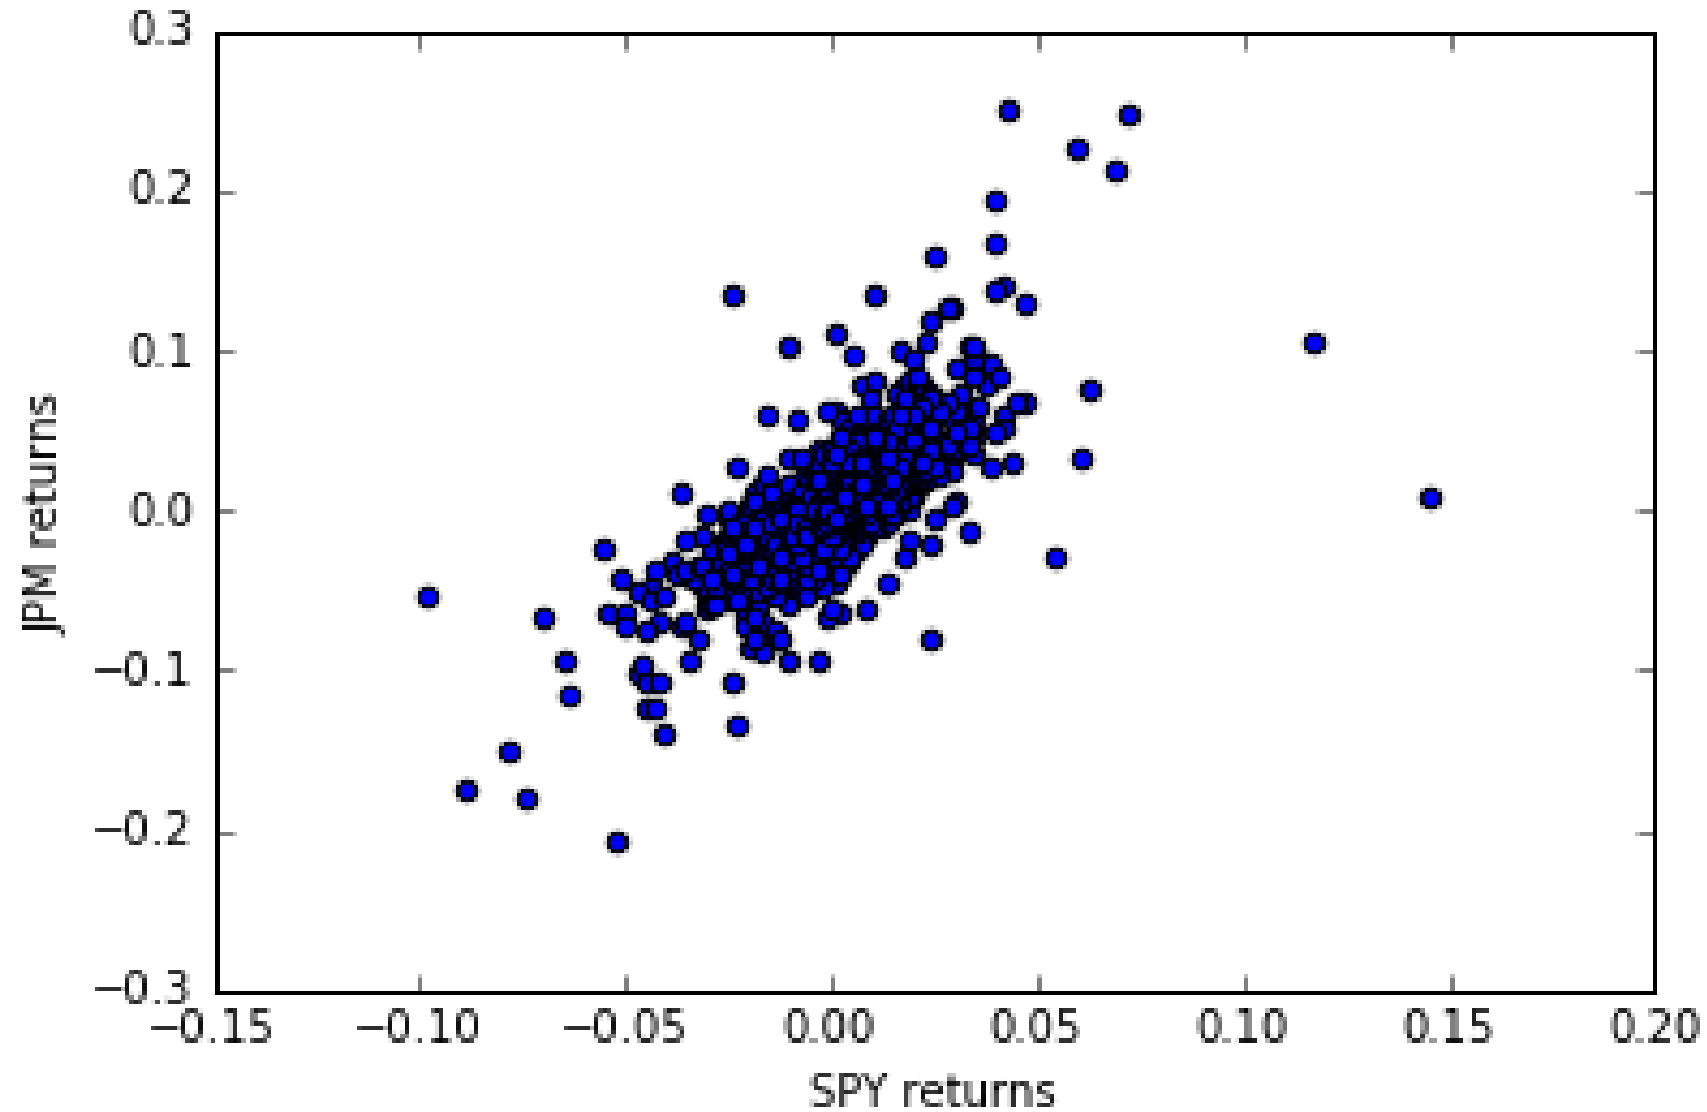
Correlation of Two Time Series

- Plot of S&P500 and JPMorgan stock



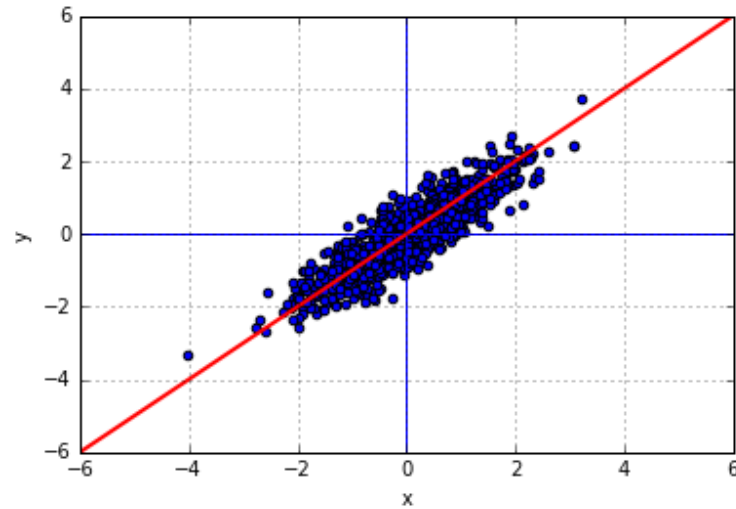
Correlation of Two Time Series

- Scatter plot of S&P500 and JP Morgan returns

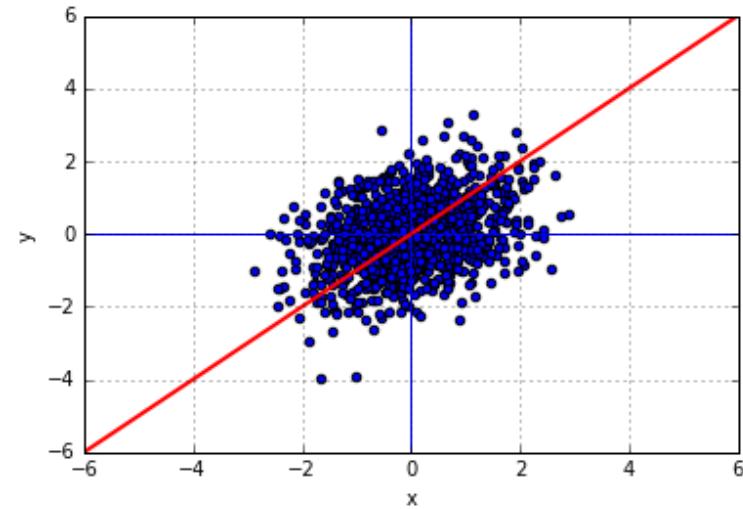


More Scatter Plots

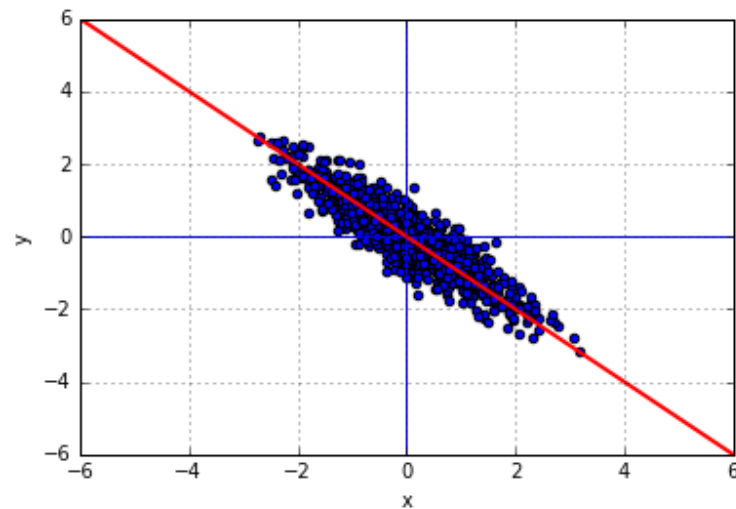
- Correlation = 0.9



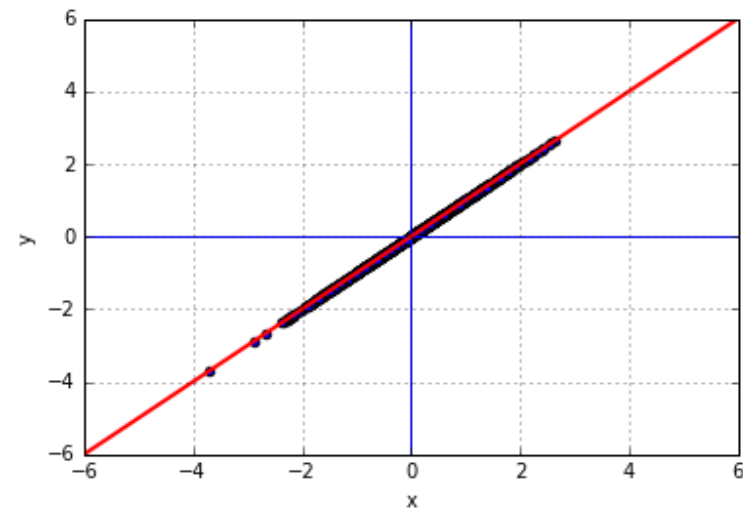
- Correlation = 0.4



- Correlation = -0.9

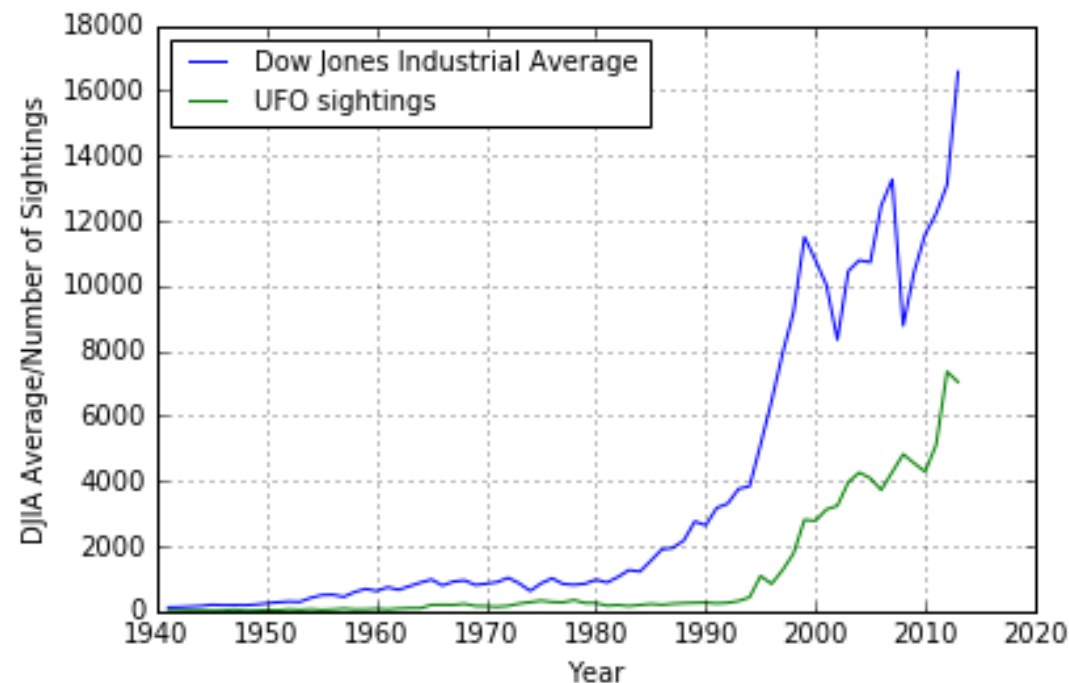


- Correlation = 1.0



Common Mistake: Correlation of Two Trending Series

- Dow Jones Industrial Average and UFO Sightings (www.nuforc.org)



- Correlation of levels: 0.94

Example: Correlation of Large Cap and Small Cap Stocks

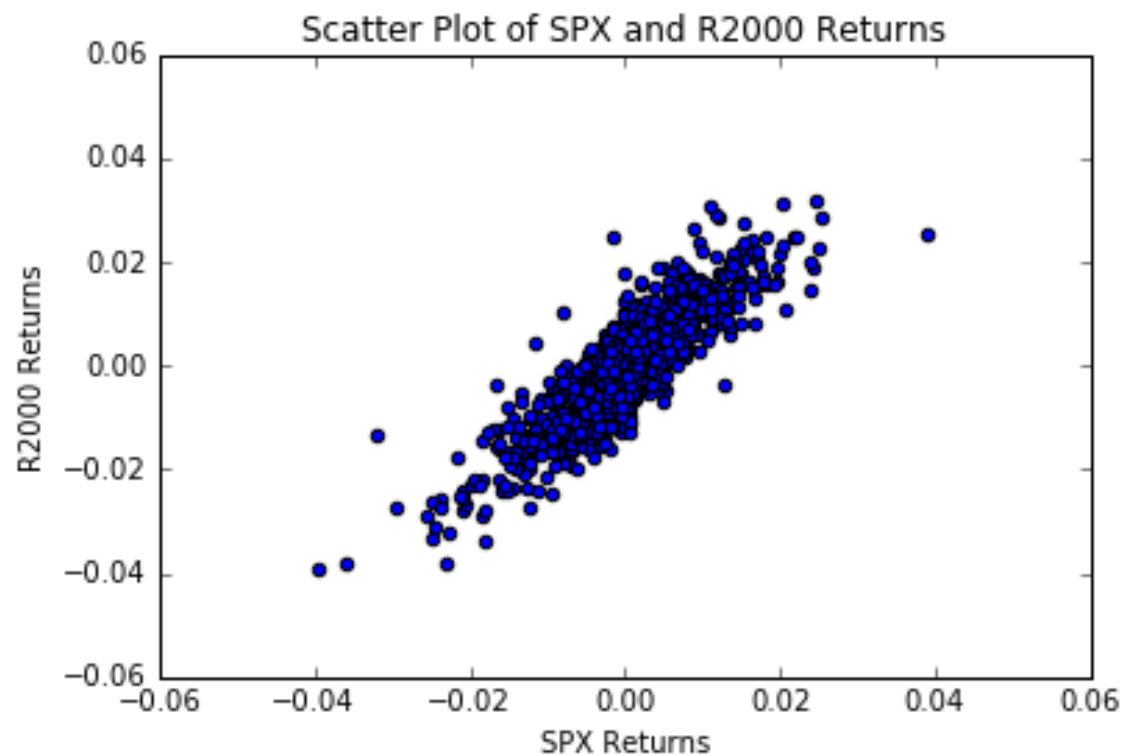
- Start with stock prices of SPX (large cap) and R2000 (small cap)
- First step: Compute percentage changes of both series

```
df['SPX_Ret'] = df['SPX_Prices'].pct_change()  
df['R2000_Ret'] = df['R2000_Prices'].pct_change()
```


Example: Correlation of Large Cap and Small Cap Stocks

- Visualize correlation with scatter plot

```
plt.scatter(df['SPX_Return'], df['R2000_Return'])  
plt.show()
```



Example: Correlation of Large Cap and Small Cap Stocks

- Use pandas correlation method for Series

```
correlation = df['SPX_Ret'].corr(df['R2000_Ret'])  
print("Correlation is: ", correlation)
```

```
Correlation is: 0.868
```

Let's practice!

INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

Simple Linear Regressions

INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON



Rob Reider

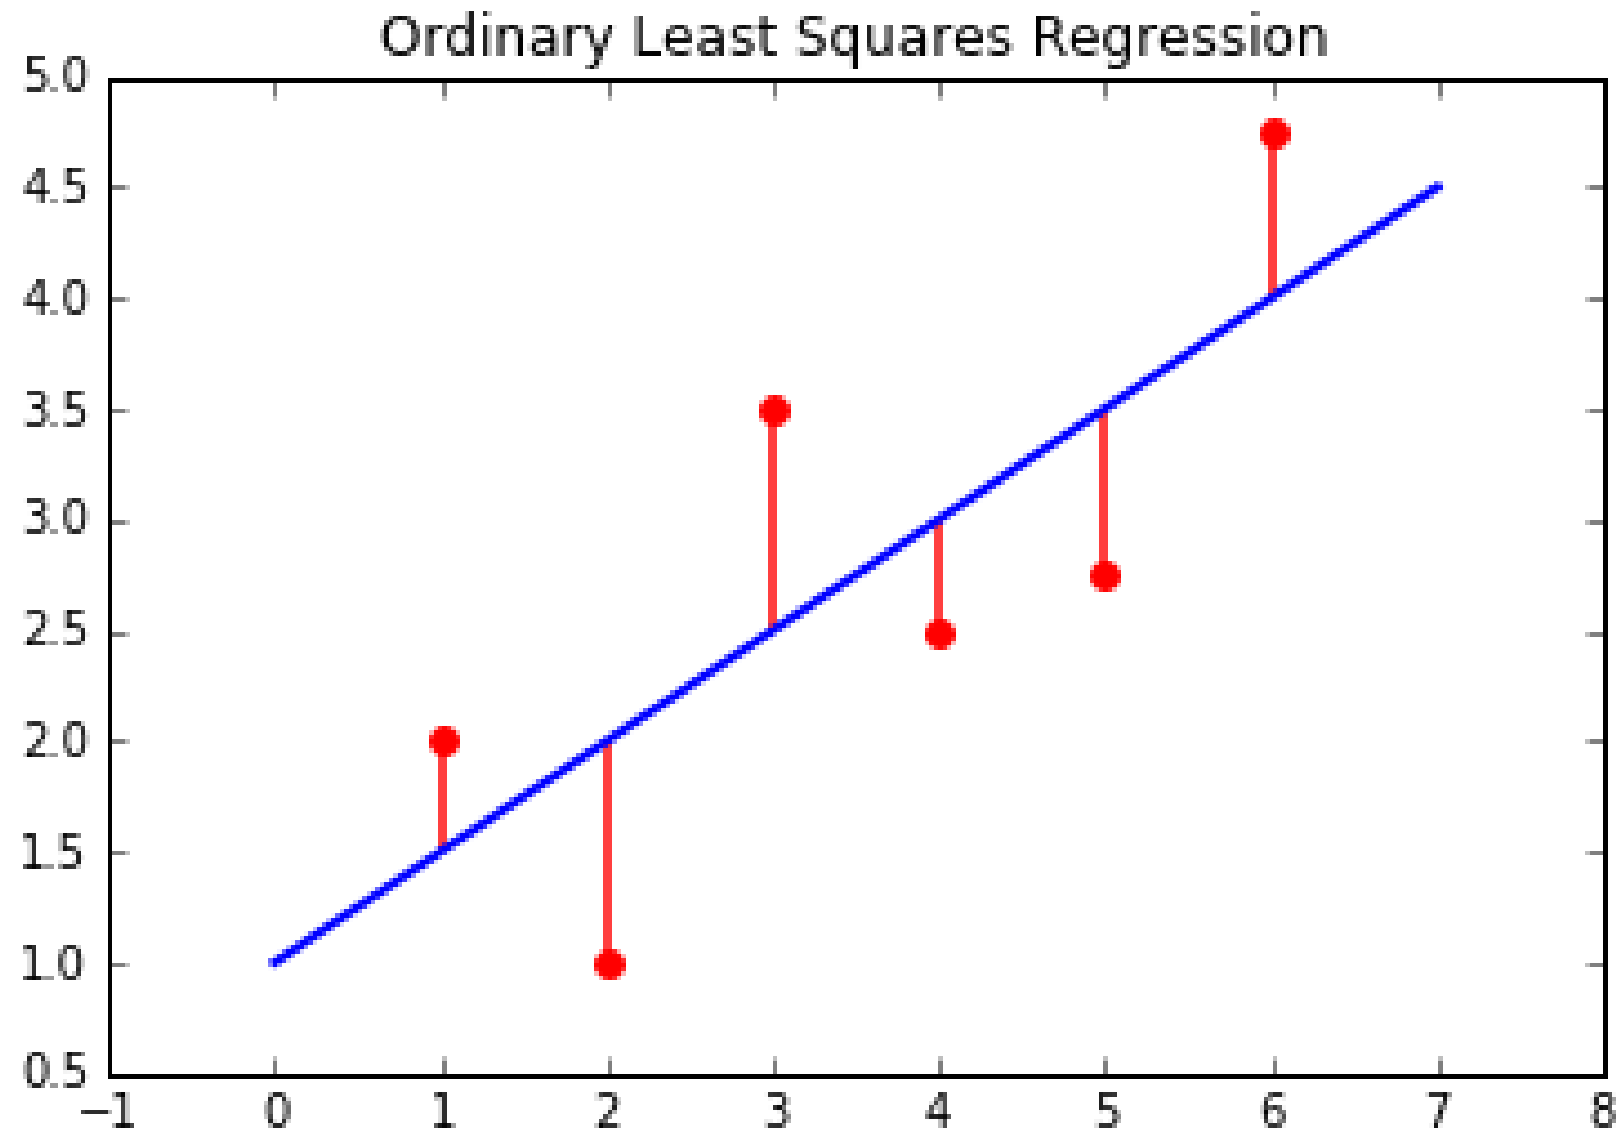
Adjunct Professor, NYU-Courant
Consultant, Quantopian

What is a Regression?

- Simple linear regression: $y_t = \alpha + \beta x_t + \epsilon_t$

What is a Regression?

- Ordinary Least Squares (OLS)



Python Packages to Perform Regressions

- In statsmodels:

```
import statsmodels.api as sm
sm.OLS(y, x).fit()
```

Warning: the order of **x** and **y** is not consistent across packages

- In numpy:

```
np.polyfit(x, y, deg=1)
```

- In pandas:

```
pd.ols(y, x)
```

- In scipy:

```
from scipy import stats
stats.linregress(x, y)
```

Example: Regression of Small Cap Returns on Large Cap

- Import the statsmodels module

```
import statsmodels.api as sm
```

- As before, compute percentage changes in both series

```
df['SPX_Ret'] = df['SPX_Prices'].pct_change()  
df['R2000_Ret'] = df['R2000_Prices'].pct_change()
```

- Add a constant to the DataFrame for the regression intercept

```
df = sm.add_constant(df)
```


Regression Example (continued)

- Notice that the first row of returns is NaN

	SPX_Price	R2000_Price	SPX_Ret	R2000_Ret
Date				
2012-11-01	1427.589966	827.849976	NaN	NaN
2012-11-02	1414.199951	814.369995	-0.009379	-0.016283

- Delete the row of NaN

```
df = df.dropna()
```

- Run the regression

```
results = sm.OLS(df['R2000_Ret'], df[['const', 'SPX_Ret']]).fit()  
print(results.summary())
```

Regression Example (continued)

- Regression output

```
=====
                        OLS Regression Results
=====
Dep. Variable:          R2000_Ret      R-squared:                0.753
Model:                  OLS           Adj. R-squared:           0.753
Method:                 Least Squares   F-statistic:              3829.
Date:                   Fri, 26 Jan 2018 Prob (F-statistic):       0.00
Time:                   13:29:55        Log-Likelihood:          4882.4
No. Observations:       1257           AIC:                     -9761.
Df Residuals:           1255           BIC:                     -9751.
Df Model:                1
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[95.0% Conf. Int.]
const	-4.964e-05	0.000	-0.353	0.724	-0.000 0.000
SPX_Ret	1.1412	0.018	61.877	0.000	1.105 1.177

```
=====
Omnibus:                61.950      Durbin-Watson:           1.991
Prob(Omnibus):           0.000      Jarque-Bera (JB):        148.100
Skew:                    0.266      Prob(JB):                6.93e-33
Kurtosis:                4.595      Cond. No.                 131.
=====
```

- Intercept in `results.params[0]`
- Slope in `results.params[1]`

Regression Example (continued)

- Regression output

```

=====
                        OLS Regression Results
=====
Dep. Variable:          R2000_Ret      R-squared:                0.753
Model:                  OLS           Adj. R-squared:           0.753
Method:                 Least Squares   F-statistic:              3829.
Date:                   Fri, 26 Jan 2018 Prob (F-statistic):       0.00
Time:                   13:29:55        Log-Likelihood:          4882.4
No. Observations:       1257           AIC:                     -9761.
Df Residuals:           1255           BIC:                     -9751.
Df Model:                1
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[95.0% Conf. Int.]
const	-4.964e-05	0.000	-0.353	0.724	-0.000 0.000
SPX_Ret	1.1412	0.018	61.877	0.000	1.105 1.177

```

=====
Omnibus:                61.950   Durbin-Watson:           1.991
Prob(Omnibus):           0.000   Jarque-Bera (JB):        148.100
Skew:                    0.266   Prob(JB):                 6.93e-33
Kurtosis:                 4.595   Cond. No.                  131.
=====

```

Relationship Between R-Squared and Correlation

- $[\text{corr}(x, y)]^2 = R^2$ (or R-squared)
- $\text{sign}(\text{corr}) = \text{sign}(\text{regression slope})$
- In last example:
 - R-Squared = 0.753
 - Slope is positive
 - $\text{correlation} = +\sqrt{0.753} = 0.868$

Let's practice!

INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

Autocorrelation

INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON



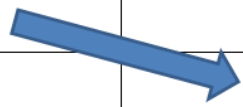
Rob Reider

Adjunct Professor, NYU-Courant
Consultant, Quantopian

What is Autocorrelation?

- Correlation of a time series with a lagged copy of itself

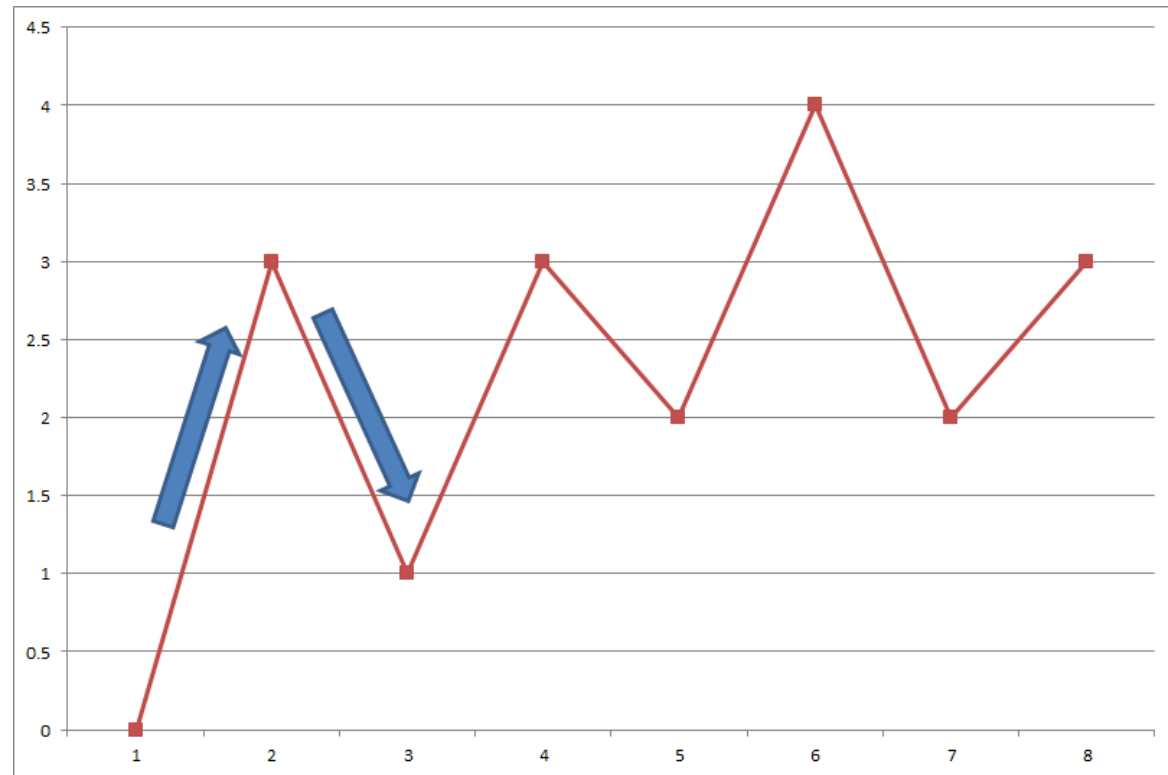
Series	Lagged Series
5	
10	5
15	10
20	15
25	20
⋮	⋮



- Lag-one autocorrelation
- Also called **serial correlation**

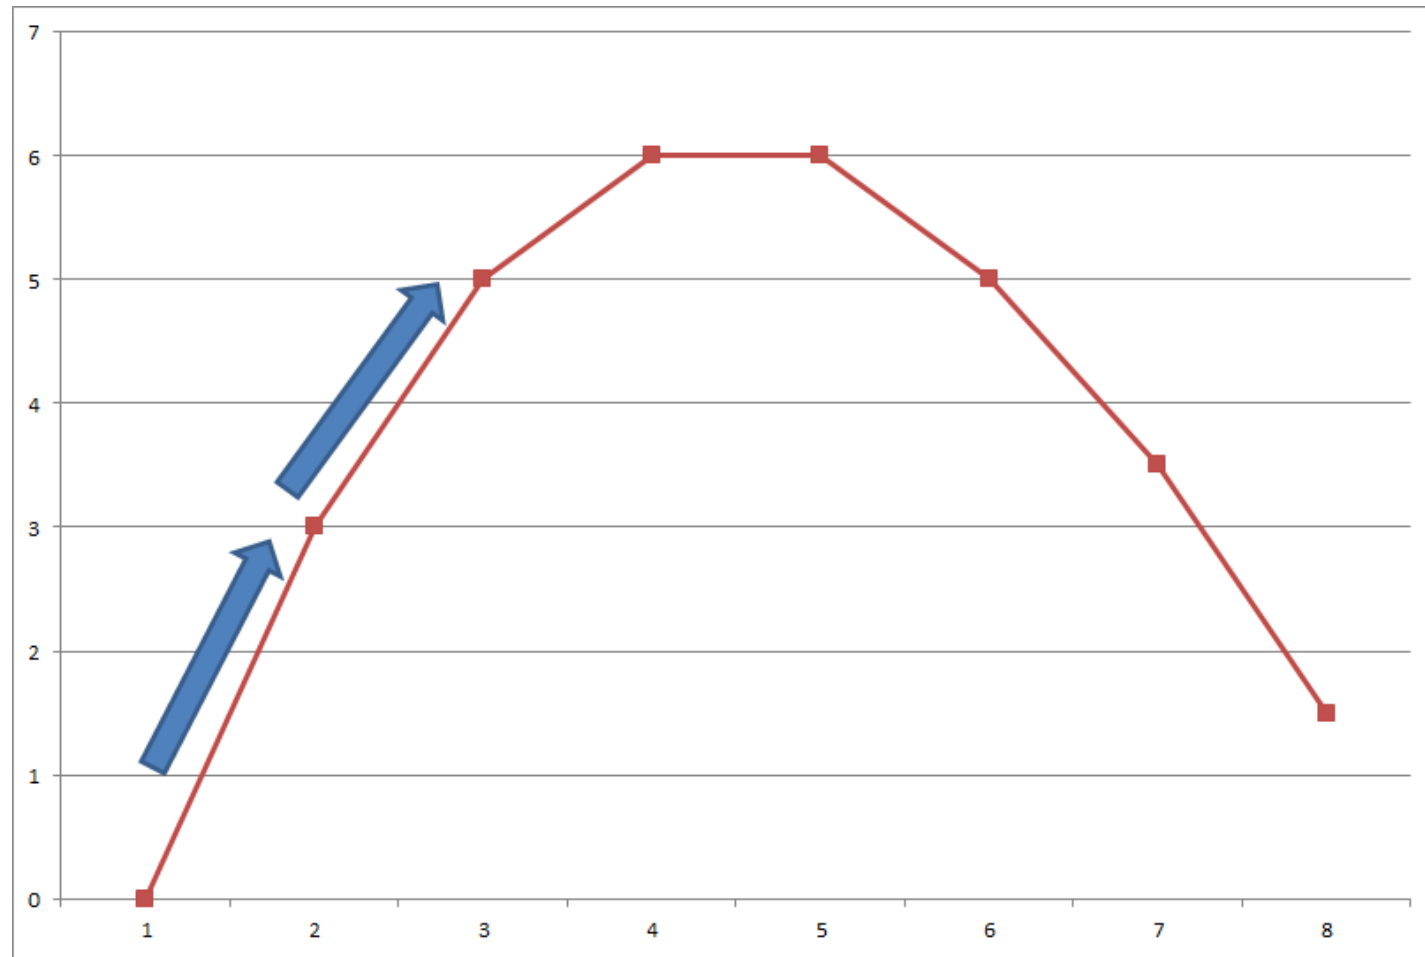
Interpretation of Autocorrelation

- Mean Reversion - Negative autocorrelation



Interpretation of Autocorrelation

- Momentum, or Trend Following - Positive autocorrelation



Traders Use Autocorrelation to Make Money

- Individual stocks
 - Historically have negative autocorrelation
 - Measured over short horizons (days)
 - Trading strategy: Buy losers and sell winners
- Commodities and currencies
 - Historically have positive autocorrelation
 - Measured over longer horizons (months)
 - Trading strategy: Buy winners and sell losers

Example of Positive Autocorrelation: Exchange Rates

- Use daily ¥/\$ exchange rates in DataFrame `df` from [FRED](#)
- Convert index to datetime

```
# Convert index to datetime
df.index = pd.to_datetime(df.index)
# Downsample from daily to monthly data
df = df.resample(rule='M', how='last')
# Compute returns from prices
df['Return'] = df['Price'].pct_change()
# Compute autocorrelation
autocorrelation = df['Return'].autocorr()
print("The autocorrelation is: ", autocorrelation)
```

```
The autocorrelation is: 0.0567
```

Let's practice!

INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON